

# Artificial Intelligence at Work

Artificial intelligence systems can support employees by searching large document collections, classifying requests and drafting answers. A useful system should not only generate fluent text. It should also make its evidence visible so that users can verify important claims. Retrieval augmented generation, usually called RAG, combines information retrieval with a language model. The retrieval component first selects relevant passages. The language model then receives those passages as context.

## Benefits and Risks

RAG can reduce hallucinations because the answer is grounded in selected source text. It does not eliminate errors. Poor document quality, unsuitable chunk sizes or an ambiguous question can still produce weak results. Organizations should evaluate both retrieval quality and answer quality. Sensitive documents also require appropriate data protection. Local embedding models keep the source text on the user's computer, while cloud-based language models may transmit retrieved passages to an external provider.

# Evaluation

A small evaluation set should contain representative questions and expected evidence. For retrieval, teams can measure whether a relevant passage appears among the top results. For generated answers, reviewers can assess correctness, completeness and whether citations actually support each claim. Comparing multiple chunk sizes and retrieval depths helps identify a reasonable configuration for the document collection.